

HACKVATAR 2026 · PERXONA CONNECT KIT

FirstRound

Your AI avatar runs the first-round interview.
It listens, adapts, and hands you a scored report — before you speak to a single candidate.

Perxona Avatar | DeepSeek LLM | Branching Dialogue Map | Auto-scored Reports

THE PROBLEM

First-round screens are broken

Every recruiter knows the pain. They just don't have a better option — until now.



Time-expensive

A recruiter spends 30–45 min on a screen call just to confirm the basics — repeating the same 6 questions a hundred times a quarter.



Forms don't work

Async questionnaires get copy-pasted LinkedIn answers. Without pressure, candidates don't reveal how they actually think.



Inconsistent

Different recruiters ask different things, follow up differently, and score differently. The data isn't comparable across candidates.

THE IDEA

A face changes everything

The avatar isn't decoration — it's the product. Candidates answering a face respond like they're in a real interview, not filling in a form.

1

Organization writes rules

Knowledge base + interview rules in plain language. DeepSeek generates a branching Dialogue Map.

2

Avatar runs the interview

The Perxona avatar speaks, listens, reacts visibly, and picks its next question from what the candidate actually said.

3

Recruiter reads the report

A scored report with verbatim quotes and a full transcript lands in the console before the recruiter opens a single CV.

UNDER THE HOOD

The Dialogue Map Language

```
// Each node is a question with branches
"greet_open": {
  "kind": "ask",
  "say": "Tell me about an LLM
        feature you shipped.",
  "motion": "greeting",
  "emotion": "curiosity",
  "branches": [
    { "when": "Names a real shipped feature",
      "to": "deep_retrieval" },
    { "when": "Vague or no real ownership",
      "to": "push_ownership" }
  ],
  "fallback": "push_ownership"
}
```

- **Adaptive routing at 1–1.5s** — the LLM reads the answer and picks the next node. Fast enough to feel like thought, not a kiosk.
- **8 named gesture slots** — the map says `probing`; one click in the motion editor changes what `probing` means for every question.
- **Emotion + intensity per node** — the avatar looks *visibly unconvinced* by a thin answer. That signal is the interview.
- **Visual editor** — recruiters can read and edit the map in the browser, no JSON required.

IMPACT

Screens that scale

FirstRound doesn't replace recruiters — it makes their judgment available at volume.

100×
Interview capacity with zero recruiter time per screen

2 min
Candidate time. Short enough that completion rates stay high

100%
Consistent questions, scoring dimensions, and verbatim evidence for every candidate

For recruiters

Write rules once. Every candidate gets the same rigorous screen. The report quotes them verbatim — no interpretation needed.

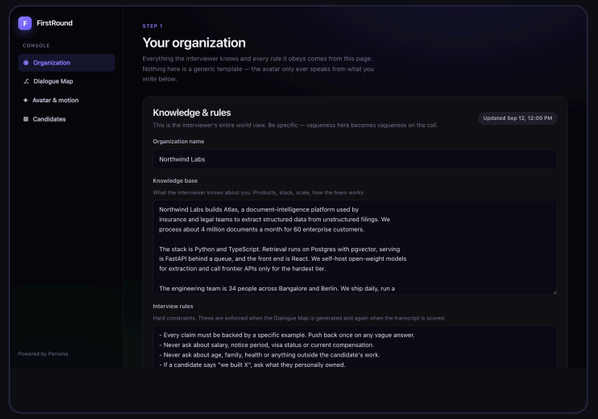
For candidates

A face, not a form. The avatar acknowledges their answers, reacts visibly, and finishes in 2 minutes — on their own schedule.

LIVE DEMO

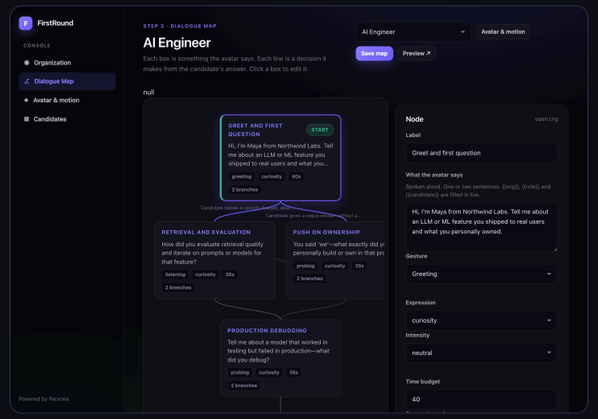
Five screens, one workflow

01 · ORGANIZATION



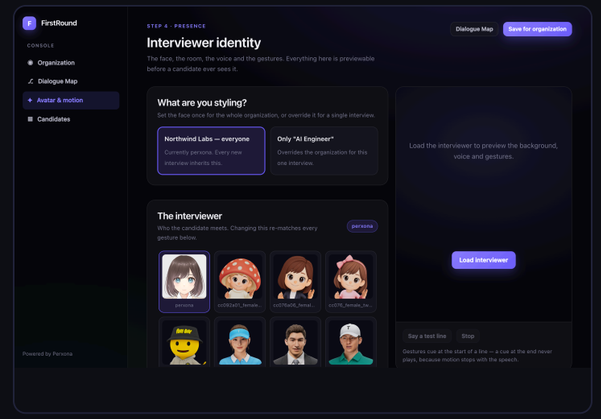
Write knowledge base + rules. Generate the Dialogue Map.

02 · DIALOGUE MAP



Visual branching editor — every question, condition, gesture.

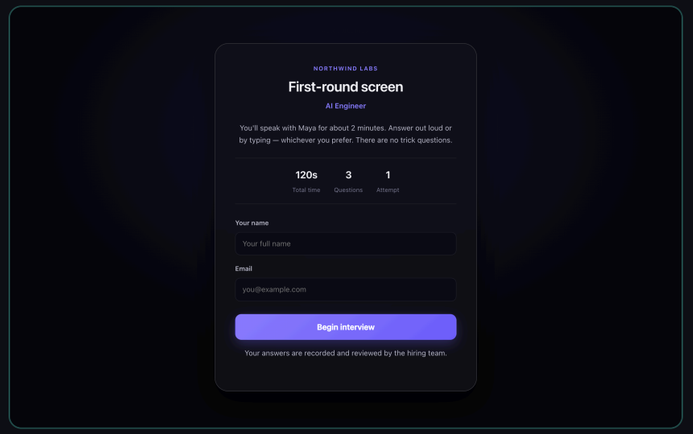
03 · AVATAR & MOTION



Pick face, scene, voice. Set 8 gesture slots with live preview.

CANDIDATE EXPERIENCE → RECRUITER READS THE EVIDENCE

04 · CANDIDATE INTERVIEW



One URL, no login. Avatar speaks, listens, reacts. 2 minutes.

05 · SCORED REPORT

